

Logistic Regression

The main tool of statistics is the Maximum-Likelihood method. Assuming the results of the experiment $\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}$, are samples from random variables $X^{(1)}, X^{(2)}, \dots, X^{(N)}$ independent and identically distributed (i.i.d.) with pdf (or pmf) $p(x)$, the joint probability of the experiment is given by

$$L(\mathbf{x}^{(1)}, \mathbf{x}^{(2)}, \dots, \mathbf{x}^{(N)}) = p(\mathbf{x}^{(1)})p(\mathbf{x}^{(2)}) \dots p(\mathbf{x}^{(N)}). \quad (1)$$

Problem 1. Consider the following random variable:

$$\begin{array}{c|c|c} Y|X=x & -1 & 1 \\ \hline P & 1-p & p \end{array}$$

Its pmf is given by $p(y) = [y = -1](1-p) + [y = 1]p$. We use the sigmoid function to transform the distance to probability:

$$P(Y = 1|X = x) = p = \sigma(\langle w, \tilde{x} \rangle) = \frac{e^{w^T \tilde{x}}}{1 + e^{w^T \tilde{x}}}.$$

If the true label is $y = 1$, then

$$p = \frac{e^{w^T \tilde{x}}}{1 + e^{w^T \tilde{x}}} = \frac{1}{1 + e^{-w^T \tilde{x}}} = \frac{1}{1 + e^{-y w^T \tilde{x}}}.$$

Prove that if $y = -1$, then $P(Y = -1|X = x) = 1 - p$ has the same expression.

Problem 2. Show that maximization of likelihood (1) leads to the loss function¹

$$Loss = \frac{1}{N} \sum_{i=1}^N \log \left(1 + e^{-y^{(i)} \langle w, \tilde{x}^{(i)} \rangle} \right).$$

Hint: Maximization of L is equivalent to maximization of $\log L$ or to minimization of $-\log L$.

¹This function is called *log-loss* or (negative) *binary cross-entropy loss*. Note that random variables $Y^{(i)}|X^{(i)} = x^{(i)}$ are not identically distributed (p depends on $x^{(i)}$), but we assume they are conditionally independent.